

Optimization in General - (i.e, not just Deep Learning)

A/Prof Richard Yi Da Xu

Yida.Xu@uts.edu.au

Wechat: aubedata

<https://github.com/roboticcam/machine-learning-notes>

University of Technology Sydney (UTS)

February 18, 2018

Gradient Descend: what is directional derivative

- ▶ **Your aim** to find:

$$\arg \min_{\mathbf{x}} (f(\mathbf{x}))$$

- ▶ How? Solve $\nabla f(\mathbf{x}_n) = 0$! But in many scenarios, this isn't easy!
- ▶ The rate of change of $f(x, y)$ in the direction of the unit vector $u = (a, b)$ is called the directional derivative $d_u f(x, y)$. The definition of the directional derivative is:


$$d_u f(x, y) = \lim_{h \rightarrow 0} \frac{f(x + ah, y + bh) - f(x, y)}{h}$$

- ▶ **Theorem** the minimum directional derivative of a differentiable function f at (x_0, y_0) is $-|\nabla f(x_0, y_0)|$ and occurs for u with the opposite direction as $\nabla f(x_0, y_0)$

Gradient Descend

Here is where **Gradient Descend** algorithm may help. The iterative algorithm looks something like:

$$\mathbf{x}_{n+1} = \mathbf{x}_n - \alpha_n \nabla f(\mathbf{x}_n), \quad n \geq 0$$



Moral of the story, you must know how to compute the objective function's **derivative**.

- Taylor expansion of $f(\mathbf{x})$ around \mathbf{x}_n in 1-D:

$$f(x_n + \Delta x) \approx f(x_n) + f'(x_n)\Delta x + \frac{1}{2}f''(x_n)\Delta x^2$$

- we need to find what is the “right” value of Δx that minimises $f(\cdot)$:

$$\frac{df(x_n + \Delta x)}{d\Delta x} = \frac{d}{d\Delta x} \left(f(x_n) + f'(x_n)\Delta x + \frac{1}{2}f''(x_n)\Delta x^2 \right) = f'(x_n) + f''(x_n)\Delta x$$

$$f'(x_n) + f''(x_n)\Delta x = 0 \implies \Delta x = \frac{-f'(x_n)}{f''(x_n)}$$

$$\begin{aligned} x_{n+1} &= x_n + \Delta x \\ &= x_n - (f''(x_n))^{-1} f'(x_n) \end{aligned}$$

- Taylor expansion of $f(\mathbf{x})$ around \mathbf{x}_n in higher dimension:

$$\implies \mathbf{x}_{n+1} = \mathbf{x}_n - \underbrace{(f''(\mathbf{x}_n))^{-1}}_{\alpha_n} \nabla f(\mathbf{x}_n)$$

- $f''(\mathbf{x}_n)$ is called Hessian matrix.

- ▶ **mean value theorem:**

if $f(x)$ is defined and **continuous** on interval $[a, b]$ and differentiable on (a, b) , then there is at least one number $c \in (a, b)$ s.t:

$$f'(c) = \frac{f(b) - f(a)}{b - a}$$

- ▶ **matrix norm:**

$$\begin{aligned}\|A\| &= \sup\{\|Ax\| : x \in K^n \text{ with } \|x\| = 1\} \\ &= \sup\left\{\frac{\|Ax\|}{\|x\|} : x \in K^n \text{ with } x \neq 0\right\}\end{aligned}$$

- ▶ when B is symmetric matrix, $\|B\| = \max\{\lambda_i(B)\}$

To prove A is symmetric matrix $\implies \|A\| = \max\{\lambda_i(A)\}$

- ▶ matrix $B = A^\top A$ is **symmetric** matrix
- ▶ **fact**: any symmetric matrix, there is an orthonormal basis of eigenvectors $\{b_i\}_{i=1}^n$, with real eigenvalues $\{\lambda_i\}_{i=1}^n$

$$Bb_i = \lambda_i b_i$$

- ▶ $B = A^\top A \implies \lambda_i$ must be non-negative real numbers, since we can write:

$$\begin{aligned} b_i^\top Bb_i &= b_i^\top \lambda_i b_i = \lambda_i \\ &= b_i^\top A^\top A b_i = (Ab_i)^\top Ab_i = \|Ab_i\|_2^2 \geq 0 \end{aligned}$$

- ▶ **unit vectors** x , i.e., $\|x\|_2 = 1$ can also be written as:

$$\left\{ x : x = \sum_{i=1}^n y_i b_i, \text{ with } \sum_{i=1}^n y_i^2 = 1 \right\}$$

this is because:

$$\begin{aligned} & (y_1 b_1^\top + y_2 b_2^\top + \cdots + y_n b_n^\top) (y_1 b_1 + y_2 b_2 + \cdots + y_n b_n) \\ &= y_1 b_1^\top (y_1 b_1 + y_2 b_2 + \cdots + y_n b_n) + \cdots + y_n b_n^\top (y_1 b_1 + y_2 b_2 + \cdots + y_n b_n) \\ &= y_1^2 + \cdots + y_n^2 \\ &= 1 \end{aligned}$$

To prove A is symmetric matrix $\implies \|A\| = \max\{\lambda_i(A)\}$

► in a same way, we can write:

$$\begin{aligned}x^\top (A^\top A)x &= x^\top Bx = \sum_{i=1}^n \lambda_i y_i^2 \\&= \left(\sum_{i=1}^n y_i b_i, \right)^\top B \left(\sum_{i=1}^n y_i b_i, \right) \\&= \sum_{i=1}^n y_i b_i^\top B y_i b_i = \sum_{i=1}^n y_i^2 b_i^\top B b_i \\&= \sum_{i=1}^n \lambda_i y_i^2\end{aligned}$$

We can now rewrite the 2-norm squared of A as

$$\begin{aligned}\|A\|_2^2 &= \max_{\{x: \|x\|=1\}} \{\|Ax\|_2^2\} = \max_{\{x: \|x\|=1\}} \{x^\top (A^\top A)x\} \\&= \max_{\{y_1, \dots, y_n \text{ s.t. } \sum y_i^2 = 1\}} \sum_{i=1}^n \lambda_i y_i^2 \\&= \max\{\lambda_i\}\end{aligned}$$

the above must occur when the y_i correspond to $\max\{\lambda_i\}$ is one, and the rest $\{y_i\}$ are zeros.

- ▶ **zero-order condition:** line above curve

$$f(\theta x + (1 - \theta)y) \leq \theta f(x) + (1 - \theta)f(y) \quad \forall x, y \in \mathbb{R}^n \quad \forall 0 \leq \theta \leq 1$$

- ▶ **first-order condition:** curve globally above tangent

$$f(y) \geq f(x) + \nabla f(x)^\top (y - x) \quad \forall x, y \in \mathbb{R}^n$$

- ▶ **second-order condition:** curve flat or curved upwards in every direction

$$0 \preceq \nabla^2 f(x) \quad \forall x \in \mathbb{R}^n$$

exercise which convex function generates a flat (constant) $\nabla^2 f(x)$

Monotonicity of gradient

- ▶ f is said to be **monotone** (non-decreasing) if $\forall (x, y), (x', y') \in \mathbb{R}^2$:

$$(x \leq x' \text{ AND } y \leq y') \implies f(x, y) \leq f(x', y')$$

think about the case of fixing one variable $y = y'$

- ▶ **exercise** does it imply monotone in both x and y ?
- ▶ then $f : \mathbb{R}^n \rightarrow \mathbb{R}^n$ is monotone mapping:

$$(f(x) - f(y))^{\top} (x - y) \geq 0$$

visualised by drawing two separate vectors $x - y$ and $f(x) - f(y)$: **both needs to be in the same quadrant**

- ▶ likewise, $\nabla f : \mathbb{R}^n \rightarrow \mathbb{R}^n$ is monotone mapping::

$$(\nabla f(x) - \nabla f(y))^{\top} (x - y) \geq 0$$

think it in terms of $\nabla^2 f$ being positive

proof of Monotonicity of gradient

- ▶ if f is differentiable and **convex**, then using **first-order condition**:

$$f(y) \geq f(x) + \nabla f(x)^\top (y - x) \qquad f(x) \geq f(y) + \nabla f(y)^\top (x - y)$$

- ▶ then the proof is:

$$-f(x) \geq -f(y) + \nabla f(x)^\top (y - x) \qquad f(x) \geq f(y) + \nabla f(y)^\top (x - y)$$

add them up:

$$\begin{aligned} 0 &\geq \nabla f(x)^\top (y - x) + \nabla f(y)^\top (x - y) \\ \nabla f(x)^\top (x - y) &\geq \nabla f(y)^\top (x - y) \\ (\nabla f(x) - \nabla f(y))^\top (x - y) &\geq 0 \end{aligned}$$

$$|f(x) - f(y)| \leq L\|x - y\|$$

this means that function f **can not** change too quickly:

- ▶ consider l_2 -regularized logistic regression, change usual notation $\theta \rightarrow x$, and $x_i \rightarrow d_i$

$$f(x) = \sum_{i=1}^n \log(1 + \exp(-y_i(x^\top d_i))) + \frac{\lambda}{2} \|x\|^2$$

- ▶ $f(x)$ is convex
- ▶ first term **is** Lipschitz continuous, second term **is not**.
- ▶ $\mu I \preceq \nabla^2 f(x) \preceq LI$ where $L = \frac{1}{4} \|A\|_2^2 + \lambda$ and $\mu = \lambda$
- ▶ gradient is Lipschitz-continuous
- ▶ function is strongly-convex

- ▶ Taylor expansion: for some z :

$$f(y) = f(x) + \nabla f(x)^\top (y - x) + \frac{1}{2}(y - x)^\top \nabla^2 f(z)(y - x)$$

- ▶ this does not look like the **usual** Taylor expansion.
- ▶ remember the mean-value theorem:

$$f'(c) = \frac{f(b) - f(a)}{b - a} \implies f(b) = f(a) + f'(c)(b - a)$$

- ▶ **mean value theorem** only gives the existence of such a point c , and not a method for how to find c

$$f(b) = [f(a) + f'(a)(b - a)] + \frac{f''(c)}{2}(b - a)^2$$

- Taylor expansion: for some z :

$$\begin{aligned} f(y) &= f(x) + \nabla f(x)^\top (y - x) + \frac{1}{2} (y - x)^\top \nabla^2 f(z) (y - x) \\ &\leq f(x) + \nabla f(x)^\top (y - x) + \frac{L}{2} \|y - x\|^2 \end{aligned}$$

- because

$$\nabla^2 f(z) \preceq LI$$

- $\nabla^2 f(z)$ is a symmetric positive definite matrix, means that

$$\begin{aligned} \nabla^2 f(z) \preceq LI &\implies \|\nabla^2 f(z)\| \leq \|LI\| \implies \|\nabla^2 f(z)\| \leq L \\ \implies \underbrace{\|\nabla^2 f(z)\|}_{\text{mean value theorem}} &= \frac{\|\nabla f(x) - \nabla f(y)\|}{\|x - y\|} \leq L \\ \implies \underbrace{\|\nabla f(x) - \nabla f(y)\|}_{\text{gradient of Lipschitz-continuous function will change at least } L} &\leq L\|x - y\| \end{aligned}$$

optimisation using Lipschitz-Continuous Gradient

- ▶ optimising the upperbound:

$$\begin{aligned}f(y) &\leq f(x) + \nabla f(x)^\top (y - x) + \frac{L}{2} \|y - x\|^2 \\ \frac{df(y)}{dy} &= \nabla f(x) + L(y - x) = 0 \\ \implies Ly &= Lx - \nabla f(x) \\ \implies y &= x - \frac{1}{L} \nabla f(x)\end{aligned}$$

- ▶ how much do we reduce? substitute $y = x - \frac{1}{L} \nabla f(x)$ into $f(y)$:
- ▶ $\nabla^2 f(x)$ is a symmetric positive definite matrix, means that

$$\begin{aligned}f(y) &\leq f(x) + \nabla f(x)^\top (y - x) + \frac{L}{2} \|y - x\|^2 \\ &= f(x) + \nabla f(x)^\top \left(x - \frac{1}{L} \nabla f(x) - x \right) + \frac{L}{2} \left\| \left(x - \frac{1}{L} \nabla f(x) \right) - x \right\|^2 \\ &= f(x) + \nabla f(x)^\top \left(-\frac{1}{L} \nabla f(x) \right) + \frac{L}{2} \left\| \frac{1}{L} \nabla f(x) \right\|^2 \\ &= f(x) - \frac{L}{2} \|\nabla f(x)\|^2 \implies f(x) - f(y) \geq \frac{L}{2} \|\nabla f(x)\|^2\end{aligned}$$

an update decreases at least $\frac{L}{2} \|\nabla f(x)\|^2$

- Taylor expansion: for some z :

$$\begin{aligned} f(y) &= f(x) + \nabla f(x)^\top (y - x) + \frac{1}{2}(y - x)^\top \nabla^2 f(z)(y - x) \\ &\leq f(x) + \nabla f(x)^\top (y - x) + \frac{L}{2}\|y - x\|^2 \end{aligned}$$

- because

$$\nabla^2 f(z) \preceq L I$$

- Traditional gradient descent approach: $\theta_{n+1} = \theta_n - \alpha_n \left(\sum_{i=1}^N x_i^\top \theta - y_i \right)$
- However, think about what if N is 1,000,000, which happens often in the BIG DATA era.
- Stochastic Gradient Descent HELPS!

A quick demo to show Stochastic Gradient Descent (1)


A simple example:

$$F(\theta) = \|\mathbf{x}^T \theta - \mathbf{y}\|^2 = \sum_{i=1}^N (x_i^T \theta - y_i)^2$$

$$\begin{aligned}\nabla F(\theta) &= 2\mathbf{x}^T (\mathbf{x}\theta - \mathbf{y}) \\ &\propto \mathbf{x}\theta - \mathbf{y} \\ &= \sum_{i=1}^N x_i^T \theta - y_i\end{aligned}$$

- ▶ Traditional gradient descent approach: $\theta_{n+1} = \theta_n - \alpha_n \left(\sum_{i=1}^N x_i^T \theta - y_i \right)$
- ▶ However, think about what if N is 1,000,000, which happens often in the BIG DATA era.
- ▶ Stochastic Gradient Descent HELPS!

A quick demo to illustrate Stochastic Gradient Descent (2)



Idea, instead of

$$\theta_{n+1} = \theta_n - \alpha_n \left(\sum_{i=1}^N x_i^T \theta - y_i \right)$$

Each iteration, we select randomly a data point pair (x_j, y_j) , and do:

$$\theta_{n+1} = \theta_n - \alpha_n (x_j^T \theta - y_j) \quad j \sim U(1, \dots, N)$$

It surprisingly works quite well in many settings. See demo

- ▶ The objective function:

$$E_D(\mathbf{w}) + \alpha E_W(\mathbf{w})$$

- ▶ Example:


$$\frac{1}{2} \sum_{n=1}^N \left(t_n - \mathbf{w}^T \phi(\mathbf{x}_n) \right)^2 + \frac{\alpha}{2} \mathbf{w}^T \mathbf{w} \implies \mathbf{w}_{\text{ML}} = \left(\alpha \mathbf{I} + \Phi^T \Phi \right)^{-1} \Phi^T \mathbf{t}$$

- ▶ A generalised example:

$$\frac{1}{2} \sum_{n=1}^N \left(t_n - \mathbf{w}^T \phi(\mathbf{x}_n) \right)^2 + \frac{\alpha}{2} \sum_{j=1}^M |w_j|^q \implies \mathbf{w}_{\text{ML}} \text{ not so easy to obtain}$$

Diagrams of ϕ_j and struggle between $E_D(\mathbf{w})$ and $\alpha E_W(\mathbf{w})$

Plot of various norm functions: q -norm $\|\mathbf{w}\|_q := \left(\sum_{i=1}^n |w_i|^q \right)^{1/q} = 1$:



minimise $E_D(\mathbf{w}) + \alpha E_W(\mathbf{w})$ becomes the “tradeoff” between the two:

